

```
In [1]: #Import Required Libraries
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [3]: #Load the dataset
Lipinski = pd.read_csv("C:\Lipinski.csv")
Lipinski.head()
```

```
Out[3]:
```

	Name	Molecular_Weight	XLogP	H-Bond_Donor_Count	H-Bond_Acceptor_Count
0	Glycyrrhizin	822.9	3.7	8	16
1	Etoposide	588.6	0.6	3	13
2	Vinblastine	811.0	3.7	3	12
3	Vincristine	825.0	2.8	3	12
4	Sorafenib	464.8	4.1	3	7

```
In [4]: #Display dataset structure including columns,dtypes,and non-null values
Lipinski.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 160 entries, 0 to 159
Data columns (total 6 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Name                  160 non-null   object
1   Molecular_Weight      160 non-null   float64
2   XLogP                 160 non-null   float64
3   H-Bond_Donor_Count    160 non-null   int64
4   H-Bond_Acceptor_Count 160 non-null   int64
5   SMILES                160 non-null   object
dtypes: float64(2), int64(2), object(2)
memory usage: 7.6+ KB
```

```
In [5]: #Show the names of all available columns
print(Lipinski.columns)
```

```
Index(['Name', 'Molecular_Weight', 'XLogP', 'H-Bond_Donor_Count',
       'H-Bond_Acceptor_Count', 'SMILES'],
      dtype='object')
```

```
In [6]: #Drop columns that are not relevant for lipinski rule analysis
Lipinski = Lipinski.drop(columns=['Name', 'SMILES'])
print(Lipinski.columns)
```

```
Index(['Molecular_Weight', 'XLogP', 'H-Bond_Donor_Count',
       'H-Bond_Acceptor_Count'],
      dtype='object')
```

```
In [7]: # Applying lipinski's rule of 5
# Returns 1 for Drug-Like (0 or 1 violations), 0 for Non-drug-Like (more than 1 vi
def Lipinski_rule(row):
    molecular_weight = row['Molecular_Weight']
    xlogp = row['XLogP']
    h_bond_donors = row['H-Bond_Donor_Count']
    h_bond_acceptors = row['H-Bond_Acceptor_Count']

    violations = 0

    if molecular_weight > 500:
        violations += 1
    if xlogp > 5:
        violations += 1
    if h_bond_donors > 5:
        violations += 1
    if h_bond_acceptors > 10:
        violations += 1

    if violations > 1:
        return 0
    else:
        return 1
```

```
In [8]: #Add 'Target' column to classify compounds for machine Learning
Lipinski['Target']=Lipinski.apply(Lipinski_rule,axis=1)
print(Lipinski)
```

	Molecular_Weight	XLogP	H-Bond_Donor_Count	H-Bond_Acceptor_Count	\
0	822.9	3.7	8	16	
1	588.6	0.6	3	13	
2	811.0	3.7	3	12	
3	825.0	2.8	3	12	
4	464.8	4.1	3	7	
..	
155	572.7	3.5	4	7	
156	588.5	1.4	9	14	
157	684.6	-1.1	7	17	
158	404.4	-0.4	6	10	
159	586.6	2.2	5	11	

	Target
0	0
1	0
2	0
3	0
4	1
..	...
155	1
156	0
157	0
158	1
159	0

[160 rows x 5 columns]

```
In [9]: # Check distribution of drug-Like (1) and non-drug-Like (0) compounds
Lipinski['Target'].value_counts()
```

```
Out[9]: Target
1      103
0       57
Name: count, dtype: int64
```

```
In [10]: # Separate features (X) and target labels (y) for machine learning
X = Lipinski[['Molecular_Weight', 'XLogP', 'H-Bond_Donor_Count',
               'H-Bond_Acceptor_Count']]
y = Lipinski['Target']
```

```
In [11]: Lipinski.shape
```

```
Out[11]: (160, 5)
```

```
In [12]: # Split the dataset into training and testing sets for model evaluation
train, test = train_test_split(Lipinski, test_size = 0.3)
print(train.shape)
print(test.shape)
```

```
(112, 5)
```

```
(48, 5)
```

```
In [13]: train_X = train[['Molecular_Weight', 'XLogP', 'H-Bond_Donor_Count',
                           'H-Bond_Acceptor_Count']]
train_y=train.Target
test_X= test[['Molecular_Weight', 'XLogP', 'H-Bond_Donor_Count',
               'H-Bond_Acceptor_Count']]
test_y =test.Target
```

```
In [14]: train_X.head(5)
```

```
Out[14]:
```

	Molecular_Weight	XLogP	H-Bond_Donor_Count	H-Bond_Acceptor_Count
132	420.6	3.8	3	6
110	435.4	-0.8	4	10
26	522.6	2.0	4	12
30	520.5	2.9	4	9
52	430.9	4.2	3	7

```
In [15]: test_X.head(5)
```

```
Out[15]:
```

	Molecular_Weight	XLogP	H-Bond_Donor_Count	H-Bond_Acceptor_Count
4	464.8	4.1	3	7
57	619.7	5.1	3	10
102	411.4	2.0	5	8
129	404.9	2.1	3	5
154	581.6	5.2	3	11

```
In [16]: train_y.head()
```

```
Out[16]: 132    1
         110    1
         26     0
         30     1
         52     1
         Name: Target, dtype: int64
```

```
In [32]: # Train Logistic Regression, SVM, Decision Tree, and KNN models and evaluate their
         model = LogisticRegression()
         model.fit(train_X,train_y)
         prediction=model.predict(test_X)
         print('The accuracy of the Logistic Regression is',accuracy_score(prediction,test_
```

The accuracy of the Logistic Regression is 0.9166666666666666

```
In [33]: from sklearn.svm import SVC
         from sklearn.tree import DecisionTreeClassifier
         from sklearn.neighbors import KNeighborsClassifier
```

```
In [34]: model = SVC(kernel= 'linear')
         model.fit(train_X,train_y)
         prediction=model.predict(test_X)
         print('The accuracy of the SVM is',accuracy_score(prediction,test_y))
```

The accuracy of the SVM is 0.8958333333333334

```
In [35]: model = DecisionTreeClassifier()
         model.fit(train_X,train_y)
         prediction=model.predict(test_X)
         print('The accuracy of the Decision Tree is',accuracy_score(prediction,test_y))
```

The accuracy of the Decision Tree is 0.9583333333333334

```
In [36]: model = KNeighborsClassifier(n_neighbors= 3)
         model.fit(train_X,train_y)
         prediction=model.predict(test_X)
         print('The accuracy of the KNN is',accuracy_score(prediction,test_y))
```

The accuracy of the KNN is 0.75

```
In [37]: #Visualization
         sns.set(style='whitegrid')
```

```
In [26]: Lipinski['Label'] = Lipinski['Target'].map({1: 'Druglike (1)', 0: 'Non-druglike (0)'})
```

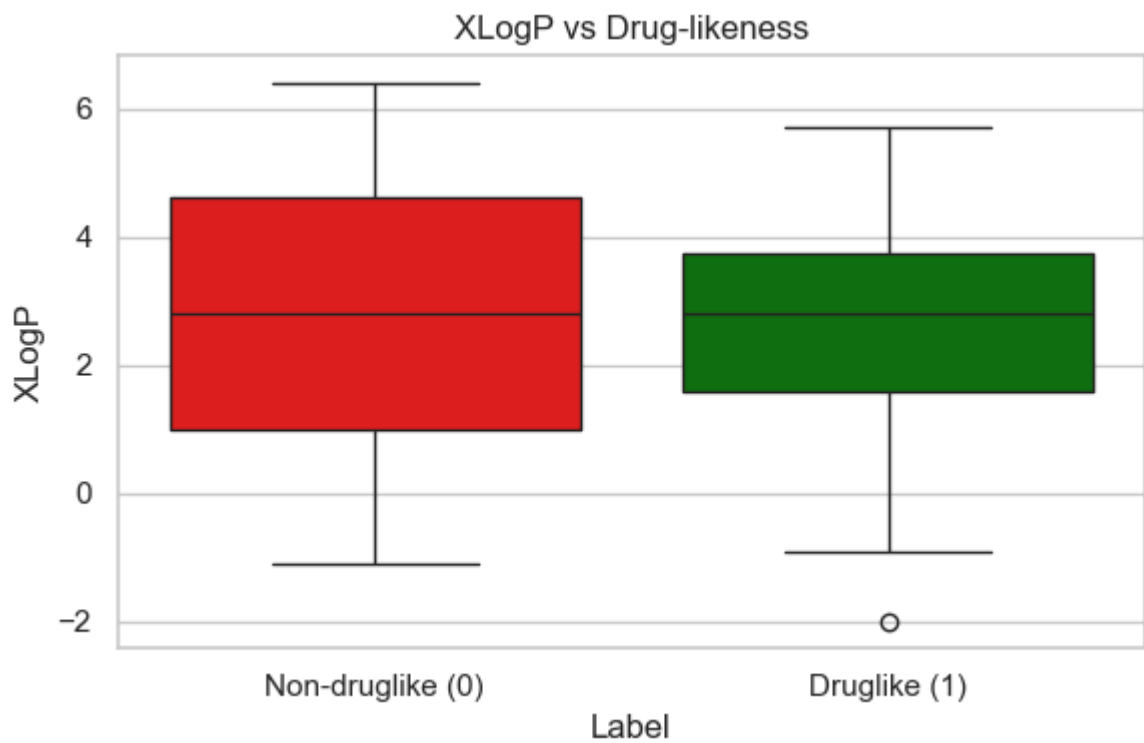
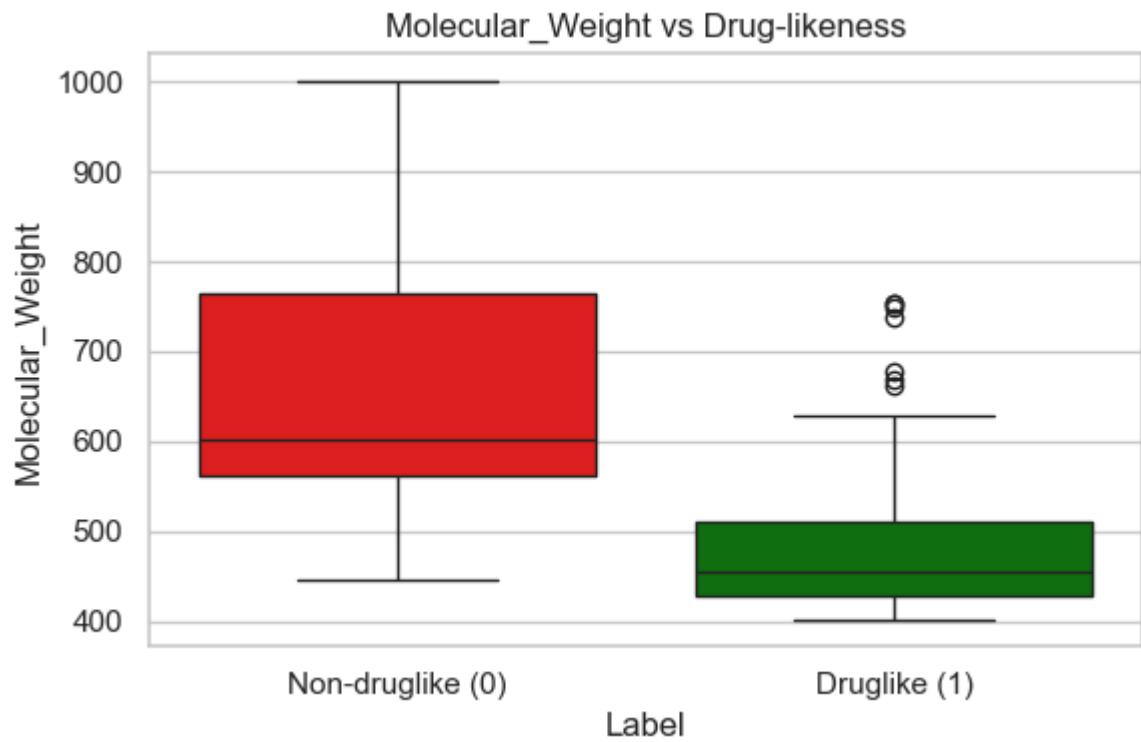
```
In [27]: palette = {'Druglike (1)': 'green', 'Non-druglike (0)': 'red'}
```

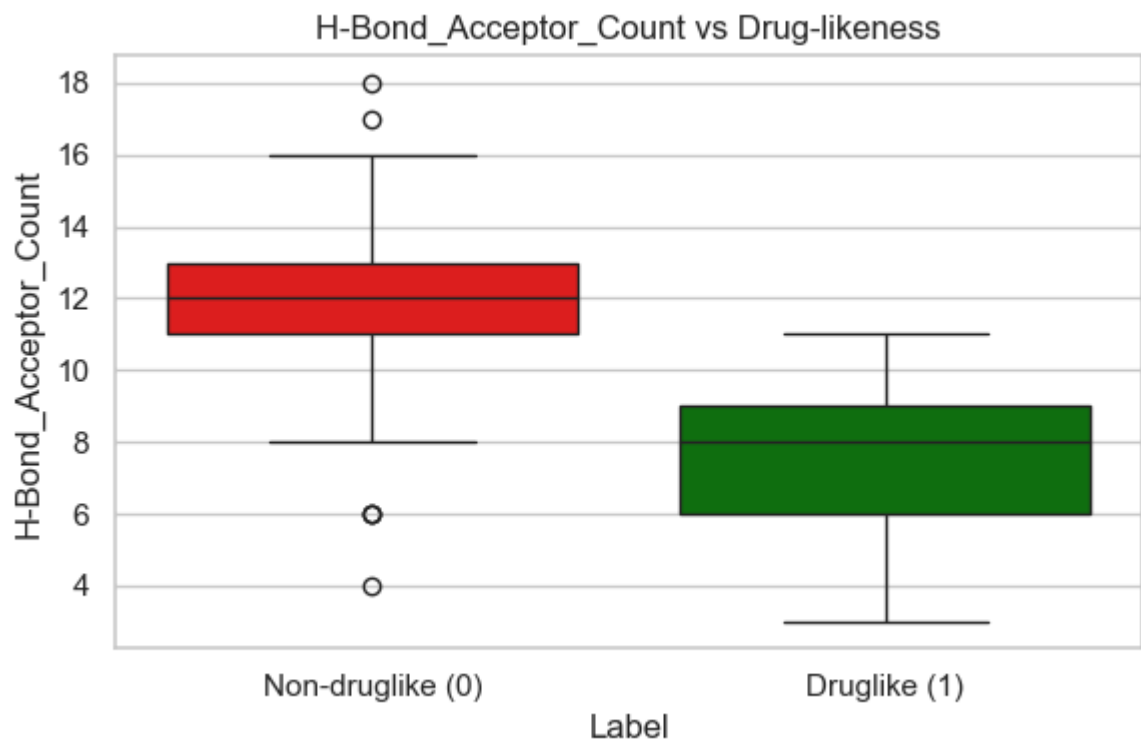
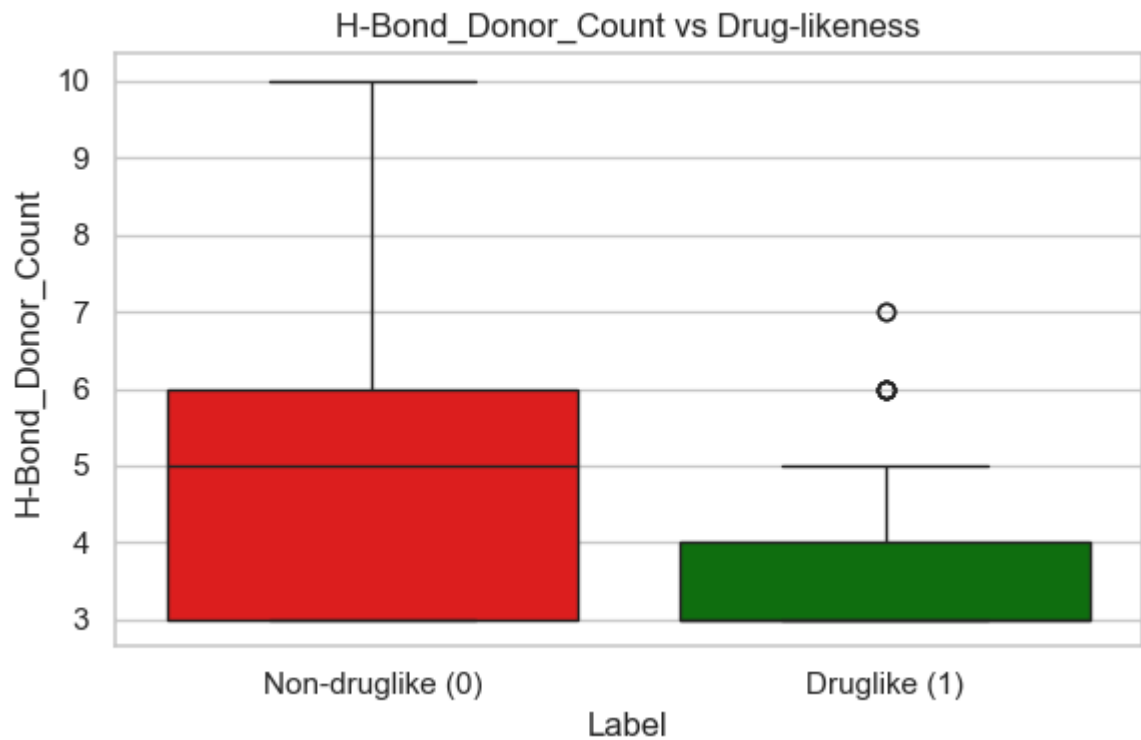
```
In [28]: # Visualize pairwise relationships between features using seaborn pairplot
         sns.pairplot(Lipinski, hue='Label', palette=palette)
         plt.suptitle("Pairwise Feature Distribution", y=1.02)
         plt.show()
```



```
In [29]: # Show distribution and outliers using a boxplot
features = ['Molecular_Weight', 'XLogP', 'H-Bond_Donor_Count', 'H-Bond_Acceptor_Count', 'Target']

for feature in features:
    plt.figure(figsize=(6, 4))
    sns.boxplot(x='Label', y=feature, data=Lipinski, hue='Label', palette=palette)
    plt.title(f"{feature} vs Drug-likeness")
    plt.tight_layout()
    plt.show()
```





```
In [31]: # Show correlation between features using a heatmap
corr = Lipinski[features + ['Target']].corr()
sns.heatmap(corr, annot=True, cmap='coolwarm', fmt=".2f")
plt.title("Feature Correlation Heatmap")
plt.show()
```

